

Class 7: Clustering and PCA


Angela Bartolo PID: A19532451

Clustering

First let's make up some data to cluster so we can get a feel for these methods and how to work with them.

We can use the 'rnorm()' function to get random numbers from a normal distribution around a given 'mean'.

```
hist(rnorm(50, mean=3, sd=4))
```




Let's get 30 points with a mean of 3 and another 30 with a mean of -3.

```
tmp <- c(rnorm(30, mean=3), rnorm(30, mean=-3))
tmp
```

```
[1] 0.1562708 1.3492176 4.0365023 4.9082740 2.4160575 3.3045219
[7] 1.8771820 2.9692954 3.4237034 2.1307541 4.1750440 0.9858158
[13] 3.8037011 2.5275276 3.1060233 1.5136745 2.2434282 1.8938065
[19] 2.8555714 2.0548467 2.4966673 3.3640423 2.9389497 1.0281483
[25] 4.6585756 4.5202011 2.5766574 4.5994736 3.0289726 3.6193110
[31] -2.3635314 -3.2447670 -2.1248421 -3.7202260 -0.9984265 -1.8487021
[37] -0.8704044 -4.4505584 -4.8619423 -1.7900690 -2.7097458 -2.5776828
[43] -3.0571633 -1.1516786 -2.1324700 -2.6908904 -4.3382307 -1.3665440
[49] -2.9914406 -3.3071907 -2.3878032 -2.3353659 -0.5748604 -2.0225735
[55] -3.1549007 -2.1071144 -4.4300802 -2.1527703 -4.5810360 -4.3145363
```

Put these two together:

```
x <- cbind(x=tmp, y=rev(tmp))
plot(x)
```



K-means clustering.

Very popular clustering method that we can use with the 'kmeans()' function in base R.

```
km <- kmeans(x, centers=2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	2.818741	-2.688585
2	-2.688585	2.818741

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 81.16257 81.16257
(between_SS / total_SS = 84.9 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Q: How many points are in each cluster?

km\$size

[1] 30 30

Q: What component of your result object details:

- Cluster size?

km\$size

[1] 30 30

- Cluster assignment/membership?

```
km$cluster
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
```

```
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


- Cluster center?

km\$centers

	x	y
1	2.818741	-2.688585
2	-2.688585	2.818741


Let's plot

```
plot(x, col=km$cluster)
points(km$centers, col= "blue", pch=15, cex =3)
```



Q Let's cluster into 3 groups or some 'x' data and make a plot.

```
km2 <- kmeans(x, centers=3)
plot(x, col=km2$cluster)
```



Hierarchical Clustering

We can use the ‘hclust()’ function for Hierarchical Clustering. Unlike ‘kmeans()’, where we could just pass in our data as input, we need to give ‘hclust()’ a “distance matrix”.

We will use the ‘dist()’ function to start with.

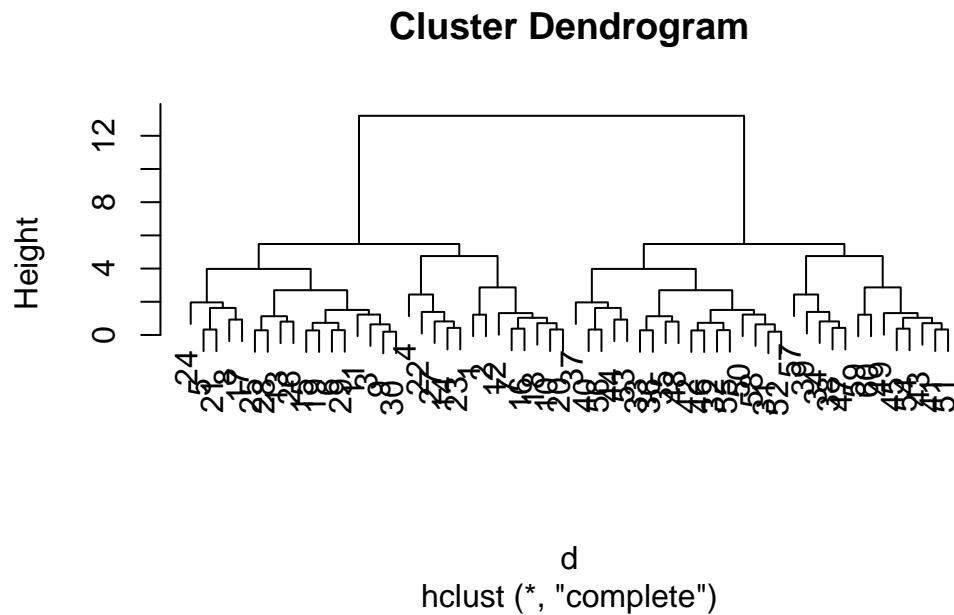
```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:

```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

```
plot(hc)
```



I can now “cut” my tree with the ‘cutree()’ to yield a cluster membership vector.

```
grps <- cutree(hc, h=8)
```


You can also tell “cutree()’ to cut where it yields”k” groups.

```
cutree(hc, k=2)
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Plot of x colored by groups

```
plot(x, col=grps)
```



Principal Component Analysis (PCA)

We will import the data provided from the UK food. 'row.names = 1' removes the first incorrect column, so the first column is counted as names and not part of the data.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1: How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17  5
```

```
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66

2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

The first column is incorrectly used in the data.

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Instead we should use 'row.names=1'.

```
x <- read.csv(url, row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139


Q2: Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

The second method is better because the first one may delete the next column if it is run again.

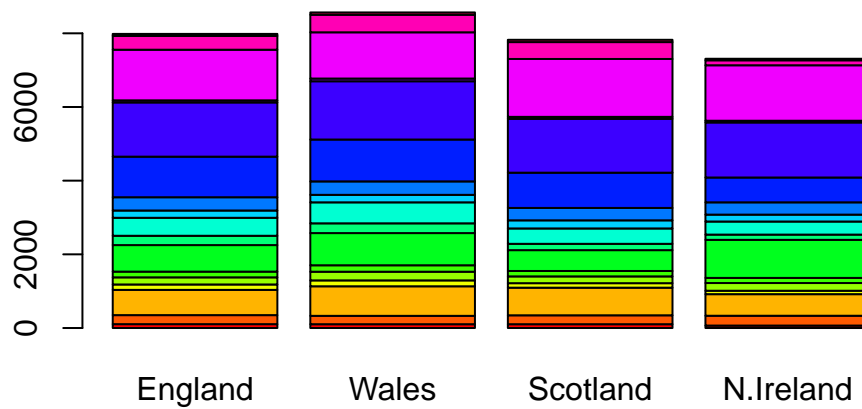
Q3: Changing what optional argument in the above barplot() function results in the following plot?

Now we will plot

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



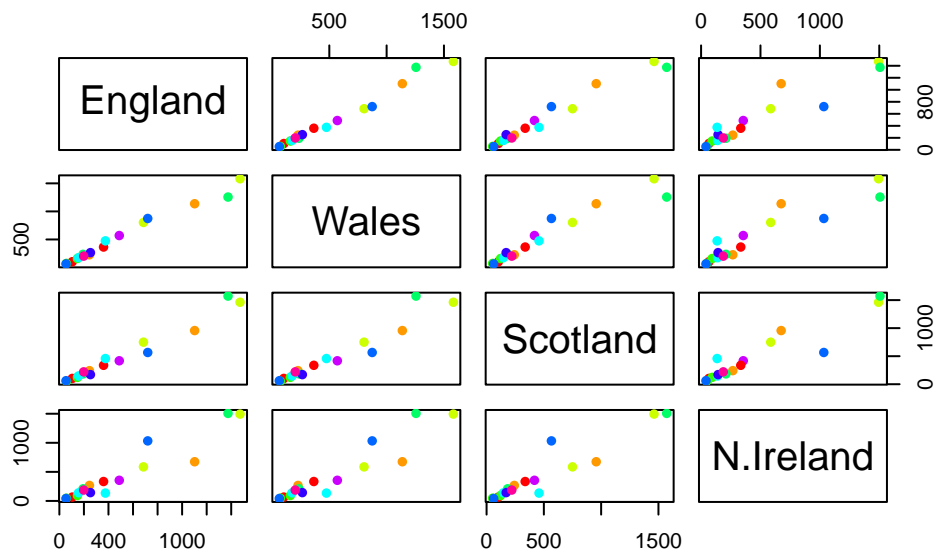
```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Using 'beside=F' places the rows of each column on top of each other instead of beside each other.

Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



This plot creates a graph to compare each country with one another. If the points are the the diagonal, the food consumption is generally the same.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

The food consumption for N. Ireland compared to other countries is not on the diagonal.

Using PCA

```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
[1] "prcomp"
```

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14
Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
paint <- c("yellow", "red", "green", "blue")
plot(pca$x[,1], pca$x[,2],
     col = paint,
     pch=16,
     xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col = paint)
```

